

# Infrastructure as Code with Terraform

Comprehensive Guide to Modern Infrastructure  
Management

*Prepared By: Rashmi Rana*  
*Corporate Trainer*

# Why Infrastructure as Code (IaC)?

---

**Infrastructure as Code (IaC)** is the practice of managing and provisioning computing infrastructure through machine-readable definition files, rather than physical hardware configuration or interactive configuration tools.

## Traditional Infrastructure Challenges

- **Manual Configuration:** Time-consuming, error-prone manual server setup and configuration
- **Configuration Drift:** Servers diverge from their intended configuration over time
- **Lack of Version Control:** No tracking of infrastructure changes or rollback capabilities
- **Inconsistent Environments:** Differences between development, staging, and production
- **Documentation Issues:** Infrastructure setup knowledge trapped in individual minds
- **Scaling Difficulties:** Manual processes don't scale with growing infrastructure needs

## IaC Benefits

**Consistency &  
Reliability**

**Version Control**

Infrastructure changes  
tracked in Git with full

Identical infrastructure across all environments, eliminating "works on my machine" issues.

history, branching, and collaboration features.

### **Automation & Speed**

Rapid provisioning and deployment of complex infrastructure in minutes, not hours or days.

### **Cost Optimization**

Programmatic resource management enables automatic scaling and cost-effective resource allocation.

## **IaC in DevOps Culture**

IaC enables the DevOps principle of treating infrastructure the same way as application code, with testing, code reviews, and continuous integration/deployment practices.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Terraform Introduction

---

**HashiCorp Terraform** is an open-source infrastructure as code software tool that provides a consistent CLI workflow to manage hundreds of cloud services. Terraform codifies cloud APIs into declarative configuration files.

## What Makes Terraform Special?

- **Multi-Cloud Support:** Works with AWS, Azure, GCP, and 1000+ other providers
- **Declarative Language:** Describe desired end state, not step-by-step instructions
- **State Management:** Tracks real-world resources and their configuration
- **Plan Before Apply:** Preview changes before making them
- **Resource Graph:** Understands dependencies and creates resources in correct order
- **Immutable Infrastructure:** Replace rather than modify existing resources

## Terraform vs Other IaC Tools

Tool	Approach	Cloud Support	Language
<b>Terraform</b>	Declarative	Multi-cloud	HCL (HashiCorp Configuration Language)
CloudFormation	Declarative	AWS only	JSON/YAML
Ansible	Procedural	Multi-cloud	YAML
Pulumi	Declarative	Multi-cloud	Python, TypeScript, Go, C#

## Terraform Architecture

Terraform uses a plugin-based architecture where providers implement resource types for specific platforms. The core Terraform engine handles state management, planning, and execution.

*Prepared By: Rashmi Rana  
Corporate Trainer*

# Terraform Providers

---

## What are Providers?

**Providers** are plugins that Terraform uses to interact with cloud platforms, SaaS providers, and other APIs. Each provider adds a set of resource types and data sources that Terraform can manage.

## Provider Configuration

```
# Configure the AWS Provider
terraform {
  required_providers {
    aws = {
      source = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}

provider "aws" {
  region = "us-west-2"
  # Optional: profile = "default"
  # Optional: access_key and secret_key
}
```

## Popular Providers

**AWS Provider**

**Azure Provider**

Manages AWS resources like EC2, S3, VPC, RDS, Lambda, and 400+ other services.

Manages Microsoft Azure resources including VMs, storage accounts, and networking.

### Google Cloud Provider

Manages GCP resources like Compute Engine, Cloud Storage, and Kubernetes Engine.

### Kubernetes Provider

Manages Kubernetes resources like deployments, services, and config maps.

## Provider Authentication

- **Environment Variables:** `AWS_ACCESS_KEY_ID`, `AWS_SECRET_ACCESS_KEY`
- **Shared Credentials File:** `~/.aws/credentials`
- **IAM Roles:** For EC2 instances or ECS tasks
- **Provider Configuration:** Direct specification in provider block

**Security Best Practice:** Never hardcode credentials in Terraform files. Use environment variables, IAM roles, or credential files instead.

# Terraform Resources

---

## Understanding Resources

**Resources** are the most important element in the Terraform language. Each resource block describes one or more infrastructure objects, such as virtual networks, compute instances, or higher-level components like DNS records.

## Resource Syntax

```
resource "resource_type" "resource_name" {  
    # Configuration arguments  
    argument1 = value1  
    argument2 = value2  
}  
  
# Example: AWS EC2 Instance  
resource "aws_instance" "web_server" {  
    ami = "ami-0c02fb55956c7d316"  
    instance_type = "t3.micro"  
    key_name = "my-key-pair"  
  
    tags = {  
        Name = "WebServer"  
        Environment = "Production"  
    }  
}
```

## Common AWS Resources



Resource Type	Purpose	Example Usage
aws_instance	EC2 virtual machines	Web servers, application servers
aws_s3_bucket	Object storage	Static websites, data storage
aws_vpc	Virtual private cloud	Network isolation
aws_security_group	Firewall rules	Network access control
aws_lb	Load balancer	Traffic distribution

## Resource Dependencies

Terraform automatically determines resource dependencies based on references between resources. You can also specify explicit dependencies using the `depends_on` argument.

```
resource "aws_security_group" "web_sg" {
  name_prefix = "web-sg"
  vpc_id = aws_vpc.main.id
}

resource "aws_instance" "web" {
  ami = "ami-0c02fb55956c7d316"
  instance_type = "t3.micro"
  vpc_security_group_ids = [aws_security_group.web_sg.id]
  # Implicit dependency on aws_security_group.web_sg
}
```

*Prepared By: Rashmi Rana  
Corporate Trainer*

# Terraform Variables

---

## Input Variables

**Input variables** serve as parameters for a Terraform module, allowing aspects of the module to be customized without altering the module's own source code.

## Variable Declaration

```
# variables.tf
variable "instance_type" {
    description = "EC2 instance type"
    type = string
    default = "t3.micro"
}

variable "environment" {
    description = "Environment name"
    type = string
    validation {
        condition = contains(["dev", "staging", "prod"],
var.environment)
        error_message = "Environment must be dev, staging, or
prod."
    }
}

variable "availability_zones" {
    description = "List of availability zones"
    type = list(string)
    default = ["us-west-2a", "us-west-2b"]
}
```

# Variable Types

## Primitive Types

string, number, bool

## Collection Types

list(type), set(type),  
map(type)

## Structural Types

object({...}),  
tuple([...])

## Special Type

any - accepts any type

# Using Variables

```
# main.tf
resource "aws_instance" "web" {
  ami = "ami-0c02fb55956c7d316"
  instance_type = var.instance_type

  tags = {
    Name = "${var.environment}-web-server"
    Environment = var.environment
  }
}
```

# Providing Variable Values

- **Command Line:** terraform apply -var="instance\_type=t3.small"
- **Variable Files:** terraform.tfvars or \*.auto.tfvars
- **Environment Variables:** TF\_VAR\_instance\_type=t3.small

- **Interactive Input:** Terraform prompts for missing variables

*Prepared By: Rashi Rana*  
*Corporate Trainer*

# Terraform Outputs

---

## Understanding Outputs

**Output values** make information about your infrastructure available on the command line and can expose information for other Terraform configurations to use.

## Output Declaration

```
# outputs.tf
output "instance_id" {
  description = "ID of the EC2 instance"
  value = aws_instance.web.id
}

output "instance_public_ip" {
  description = "Public IP address of the EC2 instance"
  value = aws_instance.web.public_ip
}

output "vpc_id" {
  description = "ID of the VPC"
  value = aws_vpc.main.id
  sensitive = false
}

output "database_password" {
  description = "Database administrator password"
  value = aws_db_instance.db.password
  sensitive = true
}
```

## Output Features

- **Description:** Human-readable description of the output value
- **Sensitive:** Mark outputs containing sensitive information
- **Depends\_on:** Explicit dependencies for outputs
- **Precondition:** Validation rules for output values

## Using Outputs

### Command Line

```
terraform output  
terraform      output  
instance_id
```

### Module Composition

Child module outputs  
become available to parent  
module

### Remote State

Access outputs from other  
Terraform configurations

### CI/CD Integration

Pass infrastructure details  
to deployment pipelines

## Complex Output Examples

```
output "load_balancer_info" {  
  description = "Load balancer information"  
  value = {  
    dns_name = aws_lb.main.dns_name  
    zone_id = aws_lb.main.zone_id  
    arn = aws_lb.main.arn  
  }  
}  
  
output "instance_ips" {
```

```
description = "List of instance IP addresses"  
value = aws_instance.web[*].private_ip  
}
```


*Prepared By: Rashmi Rana*  
*Corporate Trainer*

# Terraform Workflow

## The Core Terraform Workflow

Terraform has a simple workflow consisting of four main commands that you'll use for most operations: **init**, **plan**, **apply**, and **destroy**.

## Workflow Overview



## Command Purposes

Command	Purpose	When to Use
terraform init	Initialize working directory	First time setup, provider updates
terraform plan	Preview changes	Before applying changes, code review
terraform apply	Create/update infrastructure	Deploy changes to infrastructure



Command	Purpose	When to Use
<b>terraform destroy</b>	Remove all resources	Clean up, environment teardown

## Additional Useful Commands

- **terraform validate:** Check configuration syntax
- **terraform fmt:** Format configuration files
- **terraform show:** Display current state
- **terraform output:** Display output values
- **terraform state:** Advanced state management

**Best Practice:** Always run `terraform plan` before `terraform apply` to review changes and avoid unexpected modifications.

*Prepared By: Rashi Rana  
Corporate Trainer*

# terraform init

---

## Initialization Process

**terraform init** initializes a working directory containing Terraform configuration files. This is the first command that should be run after writing a new Terraform configuration.

## What terraform init Does

- **Downloads Providers:** Installs required provider plugins
- **Initializes Backend:** Configures state storage backend
- **Downloads Modules:** Installs child modules referenced in configuration
- **Creates Lock File:** Generates `.terraform.lock.hcl` for version consistency

## Command Examples

```
# Basic initialization
terraform init

# Initialize with backend configuration
terraform init -backend-config="bucket=my-terraform-state"

# Upgrade providers to latest versions
terraform init -upgrade

# Reconfigure backend (migrate state)
terraform init -migrate-state
```

```
# Initialize without downloading plugins (CI/CD)
terraform init -plugin-dir=/path/to/plugins
```

## Directory Structure After Init

```
my-terraform-project/
├─ main.tf
├─ variables.tf
├─ outputs.tf
├─ .terraform/ # Created by init
│   └─ providers/ # Downloaded providers
│       └─ modules/ # Downloaded modules
├─ .terraform.lock.hcl # Provider version lock
└─ terraform.tfstate # State file (after apply)
```

## Common Init Scenarios

### New Project

Run `terraform init` in directory with `.tf` files

### Provider Updates

Use `terraform init -upgrade` to update providers

### Backend Changes

Reconfigure with `terraform init -reconfigure`

### Team Collaboration

Commit `.terraform.lock.hcl` for consistent provider versions

**Important:** The .terraform directory should not be committed to version control. Add it to your .gitignore file.

*Prepared By: Rashi Rana  
Corporate Trainer*

# terraform plan

---

## Planning Infrastructure Changes

**terraform plan** creates an execution plan, showing what actions Terraform will take to reach the desired state defined in your configuration files.

## Plan Output Symbols

Symbol	Action	Description
+	Create	Resource will be created
-	Destroy	Resource will be destroyed
~	Update	Resource will be modified in-place
-/+	Replace	Resource will be destroyed and recreated
<=	Read	Data source will be read

## Plan Command Options

```
# Basic plan
terraform plan

# Save plan to file
terraform plan -out=tfplan
```

```
# Plan with variable values
terraform plan -var="instance_type=t3.small"

# Plan with variable file
terraform plan -var-file="production.tfvars"

# Target specific resources
terraform plan -target=aws_instance.web

# Detailed exit codes
terraform plan -detailed-exitcode
```

## Example Plan Output

```
Terraform will perform the following actions:

# aws_instance.web will be created
+ resource "aws_instance" "web" {
  + ami = "ami-0c02fb55956c7d316"
  + instance_type = "t3.micro"
  + id = (known after apply)
  + public_ip = (known after apply)
  + tags = {
    + "Name" = "WebServer"
  }
}

Plan: 1 to add, 0 to change, 0 to destroy.
```

## Plan Best Practices

- **Always Review:** Carefully examine plan output before applying
- **Save Plans:** Use `-out` flag for consistent apply operations
- **CI/CD Integration:** Automate plan generation in pull requests
- **Resource Targeting:** Use `-target` for selective planning



# terraform apply

---

## Applying Infrastructure Changes

**terraform apply** executes the actions proposed in a Terraform plan to create, update, or destroy infrastructure resources.

## Apply Process

- 1 **Generate Plan:** If no saved plan provided, generates execution plan
- 2 **Show Plan:** Displays planned changes for review
- 3 **Confirmation:** Prompts for approval (unless auto-approved)
- 4 **Execute Changes:** Applies changes in dependency order
- 5 **Update State:** Records current infrastructure state

## Apply Command Options

```
# Interactive apply (prompts for confirmation)
terraform apply

# Auto-approve (no confirmation prompt)
terraform apply -auto-approve

# Apply saved plan
terraform apply tfplan

# Apply with variables
terraform apply -var="environment=production"
```



```
# Target specific resources
terraform apply -target=aws_instance.web

# Parallel resource creation (default: 10)
terraform apply -parallelism=5
```

## Apply Output Example

```
aws_instance.web: Creating...
aws_instance.web: Still creating... [10s elapsed]
aws_instance.web: Still creating... [20s elapsed]
aws_instance.web: Creation complete after 22s [id=i-0123456789abcdef0]

Apply complete! Resources: 1 added, 0 changed, 0
destroyed.

Outputs:

instance_id = "i-0123456789abcdef0"
instance_public_ip = "54.123.45.67"
```

## State Management

- **State File:** terraform.tfstate tracks resource mappings
- **State Locking:** Prevents concurrent modifications
- **Remote State:** Store state in S3, Consul, or Terraform Cloud
- **State Backup:** Automatic backup before modifications

**Important:** Never manually edit the state file. Use terraform state commands for state modifications.



# terraform destroy

---

## Destroying Infrastructure

**terraform destroy** is used to destroy all resources managed by a Terraform configuration. This command is the inverse of **terraform apply**.

## Destroy Process

- 1 **Generate Destroy Plan:** Creates plan to destroy all resources
- 2 **Show Destruction Plan:** Lists resources to be destroyed
- 3 **Confirmation:** Prompts for approval with resource count
- 4 **Execute Destruction:** Destroys resources in reverse dependency order
- 5 **Clean State:** Removes destroyed resources from state

## Destroy Command Options

```
# Interactive destroy (prompts for confirmation)
terraform destroy

# Auto-approve destruction
terraform destroy -auto-approve

# Destroy specific resources
terraform destroy -target=aws_instance.web

# Destroy with variables
```

```
terraform destroy -var-file="production.tfvars"

# Plan destruction only (don't execute)
terraform plan -destroy
```

## Destroy Output Example

```
aws_instance.web: Refreshing state... [id=i-0123456789abcdef0]

Terraform will perform the following actions:

# aws_instance.web will be destroyed
- resource "aws_instance" "web" {
  - ami = "ami-0c02fb55956c7d316" -> null
  - instance_type = "t3.micro" -> null
  - id = "i-0123456789abcdef0" -> null
}

Plan: 0 to add, 0 to change, 1 to destroy.

Do you really want to destroy all resources?
  Enter a value: yes

aws_instance.web: Destroying... [id=i-0123456789abcdef0]
aws_instance.web: Still destroying... [id=i-0123456789abcdef0, 10s elapsed]
aws_instance.web: Destruction complete after 15s

Destroy complete! Resources: 1 destroyed.
```

## Destroy Safety Considerations

### Data Loss Prevention

Use `prevent_destroy` lifecycle rule for critical

### Selective Destruction

Use `-target` to destroy specific resources only

resources

### **Backup Strategy**

Backup data before  
destroying databases or  
storage

### **Environment Isolation**

Use separate state files for  
different environments

**Caution:** terraform destroy will permanently delete all managed resources. Always verify the destruction plan before confirming.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Lab: Basic Terraform Code for AWS

---

## Setting Up Your First Terraform Project

### Project Structure

```
# Create project directory
mkdir terraform-aws-lab
cd terraform-aws-lab

# Create configuration files
touch main.tf variables.tf outputs.tf
terraform.tfvars
```

### 1. Provider Configuration (main.tf)

```
terraform {
  required_version = ">= 1.0"
  required_providers {
    aws = {
      source = "hashicorp/aws"
      version = "~> 5.0"
    }
  }
}

provider "aws" {
  region = var.aws_region
}
```

### 2. Variables Definition (variables.tf)

```
variable "aws_region" {  
    description = "AWS region"  
    type = string  
    default = "us-west-2"  
}  
  
variable "environment" {  
    description = "Environment name"  
    type = string  
    default = "dev"  
}  
  
variable "instance_type" {  
    description = "EC2 instance type"  
    type = string  
    default = "t3.micro"  
}
```

### 3. Variable Values (terraform.tfvars)

```
aws_region = "us-west-2"  
environment = "development"  
instance_type = "t3.micro"
```

*Prepared By: Rashmi Rana  
Corporate Trainer*

# Lab: EC2 and S3 Resources

---

## Creating EC2 Instance and S3 Bucket

### EC2 Instance Configuration

```
# Data source for latest Amazon Linux AMI
data "aws_ami" "amazon_linux" {
  most_recent = true
  owners     = ["amazon"]

  filter {
    name = "name"
    values = ["amzn2-ami-hvm-*-x86_64-gp2"]
  }
}

# EC2 Instance
resource "aws_instance" "web_server" {
  ami = data.aws_ami.amazon_linux.id
  instance_type = var.instance_type

  tags = {
    Name = "${var.environment}-web-server"
    Environment = var.environment
    Project = "terraform-lab"
  }
}
```

### S3 Bucket Configuration

```
# S3 Bucket
resource "aws_s3_bucket" "app_bucket" {
  bucket = "${var.environment}-terraform-lab-
bucket-${random_id.bucket_suffix.hex}"
}
```



```
tags = {
  Name = "${var.environment}-app-bucket"
  Environment = var.environment
  Project = "terraform-lab"
}

# Random ID for unique bucket naming
resource "random_id" "bucket_suffix" {
  byte_length = 4
}

# S3 Bucket versioning
resource "aws_s3_bucket_versioning"
"app_bucket_versioning" {
  bucket = aws_s3_bucket.app_bucket.id
  versioning_configuration {
    status = "Enabled"
  }
}
```

## Outputs Configuration (outputs.tf)

```
output "instance_id" {
  description = "ID of the EC2 instance"
  value = aws_instance.web_server.id
}

output "instance_public_ip" {
  description = "Public IP of the EC2 instance"
  value = aws_instance.web_server.public_ip
}

output "s3_bucket_name" {
  description = "Name of the S3 bucket"
  value = aws_s3_bucket.app_bucket.bucket
}

output "s3_bucket_arn" {
  description = "ARN of the S3 bucket"
```

```
value = aws_s3_bucket.app_bucket.arn  
}
```

*Prepared By: Rashi Rana  
Corporate Trainer*

# Lab: Terraform Commands Execution

---

## Step-by-Step Execution

### 1 Initialize Terraform:

```
# Initialize the working directory
terraform init

# Expected output:
# - Downloads AWS provider
# - Creates .terraform directory
# - Creates .terraform.lock.hcl
```

### 2 Validate Configuration:

```
# Validate syntax and configuration
terraform validate

# Format configuration files
terraform fmt
```

### 3 Plan Infrastructure:

```
# Create execution plan
terraform plan

# Save plan to file
terraform plan -out=tfplan
```

```
# Expected: Shows 3 resources to add (EC2,  
S3, Random ID)
```

4

#### Apply Changes:

```
# Apply the configuration  
terraform apply  
  
# Or apply saved plan  
terraform apply tfplan  
  
# Auto-approve for automation  
terraform apply -auto-approve
```

5

#### Verify Resources:

```
# Show current state  
terraform show  
  
# Display outputs  
terraform output  
  
# Get specific output  
terraform output instance_id
```

*Prepared By: Rashi Rana  
Corporate Trainer*

# Lab: VPC and Networking

## Resources

---

### Creating VPC Infrastructure

#### VPC Configuration

```
# VPC
resource "aws_vpc" "main" {
  cidr_block = "10.0.0.0/16"
  enable_dns_hostnames = true
  enable_dns_support = true

  tags = {
    Name = "${var.environment}-vpc"
    Environment = var.environment
  }
}

# Internet Gateway
resource "aws_internet_gateway" "main" {
  vpc_id = aws_vpc.main.id

  tags = {
    Name = "${var.environment}-igw"
    Environment = var.environment
  }
}
```

#### Subnets Configuration

```
# Public Subnet
resource "aws_subnet" "public" {
```

```

count = 2
vpc_id = aws_vpc.main.id
cidr_block = "10.0.${count.index + 1}.0/24"
availability_zone =
data.aws_availability_zones.available.names[count.index]
map_public_ip_on_launch = true

tags = {
  Name = "${var.environment}-public-
subnet-${count.index + 1}"
  Environment = var.environment
  Type = "Public"
}
}

# Private Subnet
resource "aws_subnet" "private" {
  count = 2
  vpc_id = aws_vpc.main.id
  cidr_block = "10.0.${count.index + 10}.0/24"
  availability_zone =
data.aws_availability_zones.available.names[count.index]

tags = {
  Name = "${var.environment}-private-
subnet-${count.index + 1}"
  Environment = var.environment
  Type = "Private"
}
}

```

## Route Tables

```

# Public Route Table
resource "aws_route_table" "public" {
  vpc_id = aws_vpc.main.id

  route {
    cidr_block = "0.0.0.0/0"
    gateway_id = aws_internet_gateway.main.id
  }

tags = {

```

```
    Name = "${var.environment}-public-rt"
    Environment = var.environment
  }
}

# Route Table Association
resource "aws_route_table_association" "public" {
  count = length(aws_subnet.public)
  subnet_id = aws_subnet.public[count.index].id
  route_table_id = aws_route_table.public.id
}
```

*Prepared By: Rashmi Rana*  
*Corporate Trainer*

# Lab: Security Groups and Load Balancer

---

## Security and Load Balancing

### Security Groups

```
# Web Server Security Group
resource "aws_security_group" "web" {
  name_prefix = "${var.environment}-web-sg"
  vpc_id = aws_vpc.main.id

  ingress {
    description = "HTTP"
    from_port = 80
    to_port = 80
    protocol = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    description = "HTTPS"
    from_port = 443
    to_port = 443
    protocol = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  ingress {
    description = "SSH"
    from_port = 22
    to_port = 22
    protocol = "tcp"
    cidr_blocks = ["10.0.0.0/16"]
  }
}
```



```

egress {
  from_port = 0
  to_port = 0
  protocol = "-1"
  cidr_blocks = ["0.0.0.0/0"]
}

tags = {
  Name = "${var.environment}-web-sg"
  Environment = var.environment
}
}

```

## Application Load Balancer

```

# Application Load Balancer
resource "aws_lb" "main" {
  name = "${var.environment}-alb"
  internal = false
  load_balancer_type = "application"
  security_groups = [aws_security_group.alb.id]
  subnets = aws_subnet.public[*].id

  enable_deletion_protection = false

  tags = {
    Name = "${var.environment}-alb"
    Environment = var.environment
  }
}

# ALB Target Group
resource "aws_lb_target_group" "web" {
  name = "${var.environment}-web-tg"
  port = 80
  protocol = "HTTP"
  vpc_id = aws_vpc.main.id

  health_check {
    enabled = true
    healthy_threshold = 2
    interval = 30
  }
}

```

```

    matcher = "200"
    path = "/"
    port = "traffic-port"
    protocol = "HTTP"
    timeout = 5
    unhealthy_threshold = 2
}

tags = {
    Name = "${var.environment}-web-tg"
    Environment = var.environment
}
}

```

## ALB Listener

```

# ALB Listener
resource "aws_lb_listener" "web" {
    load_balancer_arn = aws_lb.main.arn
    port = "80"
    protocol = "HTTP"

    default_action {
        type = "forward"
        target_group_arn = aws_lb_target_group.web.arn
    }
}

```

*Prepared By: Rashmi Rana*  
*Corporate Trainer*

# Advanced Terraform: Modules

---

## What are Terraform Modules?

**Modules** are containers for multiple resources that are used together. A module consists of a collection of .tf files kept together in a directory.

## Module Benefits

- **Reusability:** Write once, use multiple times across projects
- **Organization:** Logical grouping of related resources
- **Encapsulation:** Hide complexity behind simple interfaces
- **Versioning:** Version control for infrastructure components
- **Testing:** Test infrastructure components in isolation

## Module Structure

```
modules/  
└─ vpc/  
    ├── main.tf # Resources  
    ├── variables.tf # Input variables  
    ├── outputs.tf # Output values  
    └─ README.md # Documentation
```

## Creating a VPC Module

```
# modules/vpc/variables.tf
variable "vpc_cidr" {
    description = "CIDR block for VPC"
    type = string
    default = "10.0.0.0/16"
}

variable "environment" {
    description = "Environment name"
    type = string
}

variable "availability_zones" {
    description = "List of availability zones"
    type = list(string)
}
```

## Using Modules

```
# main.tf
module "vpc" {
    source = "../modules/vpc"

    vpc_cidr = "10.0.0.0/16"
    environment = "production"
    availability_zones = ["us-west-2a", "us-west-2b"]
}

# Reference module outputs
resource "aws_instance" "web" {
    ami = "ami-0c02fb55956c7d316"
    instance_type = "t3.micro"
    subnet_id = module.vpc.public_subnet_ids[0]
}
```

## Module Sources

- **Local Paths:** `source = "../modules/vpc"`

- **Git Repositories:** source =  
"git::https://github.com/user/repo.git"
- **Terraform Registry:** source = "terraform-aws-modules/vpc/aws"
- **HTTP URLs:** source =  
"https://example.com/module.zip"

*Prepared By: Rashi Rana  
Corporate Trainer*

# State Files and Remote Backend

---

## Terraform State

**Terraform state** is used to map real world resources to your configuration, keep track of metadata, and improve performance for large infrastructures.

## State File Challenges

- **Team Collaboration:** Multiple developers need access to same state
- **State Locking:** Prevent concurrent modifications
- **Security:** State files may contain sensitive information
- **Backup:** Prevent state file loss
- **Versioning:** Track state changes over time

## Remote Backend Configuration

```
# backend.tf
terraform {
  backend "s3" {
    bucket = "my-terraform-state-bucket"
    key    = "prod/terraform.tfstate"
    region = "us-west-2"
    encrypt = true
    dynamodb_table = "terraform-state-lock"
  }
}
```

## S3 Backend Setup

```
# Create S3 bucket for state
resource "aws_s3_bucket" "terraform_state" {
  bucket = "my-terraform-state-bucket"
}

resource "aws_s3_bucket_versioning" "terraform_state" {
  bucket = aws_s3_bucket.terraform_state.id
  versioning_configuration {
    status = "Enabled"
  }
}

resource
"aws_s3_bucket_server_side_encryption_configuration"
"terraform_state" {
  bucket = aws_s3_bucket.terraform_state.id

  rule {
    apply_server_side_encryption_by_default {
      sse_algorithm = "AES256"
    }
  }
}

# DynamoDB table for state locking
resource "aws_dynamodb_table" "terraform_state_lock" {
  name = "terraform-state-lock"
  billing_mode = "PAY_PER_REQUEST"
  hash_key = "LockID"

  attribute {
    name = "LockID"
    type = "S"
  }
}
```

## State Management Commands

```
# List resources in state
terraform state list

# Show specific resource
terraform state show aws_instance.web

# Move resource in state
terraform state mv aws_instance.old aws_instance.new

# Remove resource from state
terraform state rm aws_instance.web

# Import existing resource
terraform import aws_instance.web i-1234567890abcdef0
```

*Prepared By: Rashmi Rana*  
*Corporate Trainer*



# Data Sources and Provisioners

---

## Data Sources

**Data sources** allow Terraform to use information defined outside of Terraform, defined by another separate Terraform configuration, or modified by functions.

## Common Data Sources

```
# Get latest Amazon Linux AMI
data "aws_ami" "amazon_linux" {
  most_recent = true
  owners     = ["amazon"]

  filter {
    name = "name"
    values = ["amzn2-ami-hvm-*-x86_64-gp2"]
  }
}

# Get availability zones
data "aws_availability_zones" "available" {
  state = "available"
}

# Get current AWS caller identity
data "aws_caller_identity" "current" {}

# Get existing VPC
data "aws_vpc" "existing" {
  filter {
    name = "tag:Name"
    values = ["production-vpc"]
  }
}
```

```
}  
}
```

## Using Data Sources

```
resource "aws_instance" "web" {  
  ami = data.aws_ami.amazon_linux.id  
  instance_type = "t3.micro"  
  availability_zone =  
    data.aws_availability_zones.available.names[0]  
  
  tags = {  
    Name = "web-server"  
    Owner = data.aws_caller_identity.current.user_id  
  }  
}
```

## Provisioners

**Provisioners** can be used to model specific actions on the local machine or on a remote machine in order to prepare servers or other infrastructure objects for service.

## File and Remote-exec Provisioners

```
resource "aws_instance" "web" {  
  ami = data.aws_ami.amazon_linux.id  
  instance_type = "t3.micro"  
  key_name = "my-key"  
  
  # File provisioner  
  provisioner "file" {  
    source = "app.conf"  
    destination = "/tmp/app.conf"  
  
    connection {
```

```

        type = "ssh"
        user = "ec2-user"
        private_key = file("~/ssh/id_rsa")
        host = self.public_ip
    }
}

# Remote-exec provisioner
provisioner "remote-exec" {
    inline = [
        "sudo yum update -y",
        "sudo yum install -y httpd",
        "sudo systemctl start httpd",
        "sudo systemctl enable httpd"
    ]

    connection {
        type = "ssh"
        user = "ec2-user"
        private_key = file("~/ssh/id_rsa")
        host = self.public_ip
    }
}
}

```

**Best Practice:** Use provisioners as a last resort. Consider using user\_data, cloud-init, or configuration management tools like Ansible instead.

*Prepared By: Rashmi Rana  
Corporate Trainer*

# Hands-on: 2-Tier Architecture (Web + DB)

---

## Architecture Overview

We'll deploy a complete 2-tier architecture with web servers in public subnets and a database in private subnets, including load balancing and security groups.

## Architecture Components

### Network Layer

VPC, Public/Private Subnets, Internet Gateway, NAT Gateway, Route Tables

### Web Tier

Application Load Balancer, Auto Scaling Group, EC2 Instances in Public Subnets

### Database Tier

RDS MySQL Database in Private Subnets with Multi-AZ deployment

### Security

Security Groups, NACLs, IAM Roles, Parameter Store for secrets

## Project Structure

```
2-tier-architecture/  
├─ main.tf # Main configuration  
├─ variables.tf # Input variables  
├─ outputs.tf # Output values  
├─ terraform.tfvars # Variable values  
├─ modules/  
│   ├─ networking/ # VPC, subnets, gateways  
│   ├─ security/ # Security groups  
│   ├─ compute/ # EC2, ALB, ASG  
│   └─ database/ # RDS configuration  
└─ userdata/  
    └─ web_server.sh # EC2 initialization script
```

## Key Features

- **High Availability:** Multi-AZ deployment across 2 availability zones
- **Scalability:** Auto Scaling Group for web tier
- **Security:** Database in private subnets, security groups
- **Load Balancing:** Application Load Balancer with health checks
- **Monitoring:** CloudWatch integration

*Prepared By: Rashmi Rana  
Corporate Trainer*

# 2-Tier Architecture Implementation

---

## Database Configuration

### RDS Subnet Group

```
resource "aws_db_subnet_group" "main" {
  name = "${var.environment}-db-subnet-group"
  subnet_ids = var.private_subnet_ids

  tags = {
    Name = "${var.environment}-db-subnet-group"
    Environment = var.environment
  }
}
```

### RDS Instance

```
resource "aws_db_instance" "main" {
  identifier = "${var.environment}-database"

  engine = "mysql"
  engine_version = "8.0"
  instance_class = "db.t3.micro"

  allocated_storage = 20
  max_allocated_storage = 100
  storage_type = "gp2"
  storage_encrypted = true

  db_name = var.database_name
  username = var.database_username
  password = var.database_password

  vpc_security_group_ids =
```

```

[aws_security_group.database.id]
  db_subnet_group_name =
aws_db_subnet_group.main.name

  backup_retention_period = 7
  backup_window = "03:00-04:00"
  maintenance_window = "sun:04:00-sun:05:00"

  multi_az = true
  publicly_accessible = false
  skip_final_snapshot = true

  tags = {
    Name = "${var.environment}-database"
    Environment = var.environment
  }
}

```

## Auto Scaling Group

```

resource "aws_launch_template" "web" {
  name_prefix = "${var.environment}-web-"
  image_id = data.aws_ami.amazon_linux.id
  instance_type = var.instance_type
  key_name = var.key_name

  vpc_security_group_ids =
[aws_security_group.web.id]

  user_data =
base64encode(templatefile("${path.module}/userdata/web_
{
  db_endpoint = aws_db_instance.main.endpoint
  db_name = var.database_name
}))

  tag_specifications {
    resource_type = "instance"
    tags = {
      Name = "${var.environment}-web-server"
      Environment = var.environment
    }
  }
}

```

```
}

resource "aws_autoscaling_group" "web" {
  name = "${var.environment}-web-asg"
  vpc_zone_identifier = var.public_subnet_ids
  target_group_arns = [aws_lb_target_group.web.arn]
  health_check_type = "ELB"

  min_size = 2
  max_size = 6
  desired_capacity = 2

  launch_template {
    id = aws_launch_template.web.id
    version = "$Latest"
  }

  tag {
    key = "Name"
    value = "${var.environment}-web-asg"
    propagate_at_launch = false
  }
}
```

*Prepared By: Rashi Rana*  
*Corporate Trainer*



# Summary and Best Practices

---

## Key Concepts Covered

- Infrastructure as Code principles and benefits
- Terraform fundamentals: providers, resources, variables, outputs
- Terraform workflow: init, plan, apply, destroy
- Advanced features: modules, state management, data sources
- Practical implementation of AWS infrastructure
- 2-tier architecture deployment with best practices

## Terraform Best Practices

### Code Organization

Use modules, consistent naming, proper file structure, and documentation

### State Management

Remote backend, state locking, separate environments, regular backups

### Security

No hardcoded secrets, least privilege access, encrypted

### Version Control

Git workflows, .gitignore, lock files, semantic

state, secure backends

versioning for modules

## Production Considerations

- **Environment Separation:** Separate state files and configurations for dev/staging/prod
- **CI/CD Integration:** Automated planning and applying in pipelines
- **Monitoring:** CloudWatch, logging, and alerting for infrastructure changes
- **Disaster Recovery:** Multi-region deployments and backup strategies
- **Cost Optimization:** Resource tagging, right-sizing, and automated cleanup

## Next Steps

1. **Advanced Terraform:** Workspaces, dynamic blocks, `for_each` loops
2. **Testing:** Terratest, kitchen-terraform, policy as code
3. **Multi-Cloud:** Azure, GCP provider usage
4. **Terraform Cloud:** Remote execution, policy enforcement
5. **Integration:** Ansible, Kubernetes, service mesh

**Remember:** Infrastructure as Code is not just about tools—it's about culture, collaboration, and treating infrastructure with the same discipline as application code.